

updated

January 31, 2026

```
[ ]: # CS 8316 - BIG DATA PROGRAMMING
## Group Assignment Report

### Group Members
1. Elikana Jeremiah Dakho
2. Valerian Onesmo

---
```

```
[111]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[112]: # Load the dataset
df = pd.read_csv("/home/karera/Desktop/SEMISTER11/big data/Churndata(1).csv")
```

```
[114]: df.head()
```

```
[114]:    customerID  gender  SeniorCitizen Partner Dependents  tenure PhoneService \
0    7590-VHVEG  Female          0      Yes        No         1        No
1    5575-GNVDE    Male          0       No        No        34      Yes
2    3668-QPYBK    Male          0       No        No         2      Yes
3    7795-CFOCW    Male          0       No        No        45        No
4    9237-HQKITU  Female          0       No        No         2      Yes

      MultipleLines InternetService OnlineSecurity ... DeviceProtection \
0  No phone service             DSL           No   ...
1            No                 DSL           Yes   ...
2            No                 DSL           Yes   ...
3  No phone service             DSL           Yes   ...
4            No     Fiber optic           No   ...

TechSupport StreamingTV StreamingMovies          Contract PaperlessBilling \
0            No          No           No Month-to-month        Yes
1            No          No           No     One year        No
2            No          No           No Month-to-month      Yes
3            Yes         No           No     One year        No
```

4	No	No	No	Month-to-month	Yes
		PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check		29.85	29.85	No
1	Mailed check		56.95	1889.5	No
2	Mailed check		53.85	108.15	Yes
3	Bank transfer (automatic)		42.30	1840.75	No
4	Electronic check		70.70	151.65	Yes

[5 rows x 21 columns]

[115]: df.tail()

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
7038	6840-RESVB	Male		0	Yes	Yes	24
7039	2234-XADUH	Female		0	Yes	Yes	72
7040	4801-JZAZL	Female		0	Yes	Yes	11
7041	8361-LTMKD	Male		1	Yes	No	4
7042	3186-AJIEK	Male		0	No	No	66
	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\	
7038	Yes	Yes	DSL	Yes	...		
7039	Yes	Yes	Fiber optic	No	...		
7040	No	No phone service	DSL	Yes	...		
7041	Yes	Yes	Fiber optic	No	...		
7042	Yes	No	Fiber optic	Yes	...		
	DeviceProtection	TechSupport	StreamingTV	StreamingMovies		Contract	\
7038	Yes	Yes	Yes	Yes		One year	
7039	Yes	No	Yes	Yes		One year	
7040	No	No	No	No	Month-to-month		
7041	No	No	No	No	Month-to-month		
7042	Yes	Yes	Yes	Yes	Two year		
	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	\		
7038	Yes	Mailed check	84.80	1990.5			
7039	Yes	Credit card (automatic)	103.20	7362.9			
7040	Yes	Electronic check	29.60	346.45			
7041	Yes	Mailed check	74.40	306.6			
7042	Yes	Bank transfer (automatic)	105.65	6844.5			
	Churn						
7038	No						
7039	No						
7040	No						
7041	Yes						
7042	No						

```
[5 rows x 21 columns]
```

```
[116]: df.shape
```

```
[116]: (7043, 21)
```

```
[117]: df.columns
```

```
[117]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
       'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
       'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
       'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
       'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

```
[118]: df.dtypes
```

```
[118]: customerID          object
       gender            object
       SeniorCitizen     int64
       Partner           object
       Dependents        object
       tenure            int64
       PhoneService      object
       MultipleLines     object
       InternetService   object
       OnlineSecurity    object
       OnlineBackup       object
       DeviceProtection  object
       TechSupport        object
       StreamingTV        object
       StreamingMovies    object
       Contract           object
       PaperlessBilling  object
       PaymentMethod      object
       MonthlyCharges    float64
       TotalCharges       object
       Churn              object
       dtype: object
```

```
[119]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   customerID      7043 non-null    object 
 1   gender          7043 non-null    object 
 2   SeniorCitizen   7043 non-null    int64  
 3   Partner         7043 non-null    object 
 4   Dependents      7043 non-null    object 
 5   tenure          7043 non-null    int64  
 6   PhoneService    7043 non-null    object 
 7   MultipleLines   7043 non-null    object 
 8   InternetService 7043 non-null    object 
 9   OnlineSecurity  7043 non-null    object 
 10  OnlineBackup    7043 non-null    object 
 11  DeviceProtection 7043 non-null    object 
 12  TechSupport     7043 non-null    object 
 13  StreamingTV     7043 non-null    object 
 14  StreamingMovies 7043 non-null    object 
 15  Contract        7043 non-null    object 
 16  PaperlessBilling 7043 non-null    object 
 17  PaymentMethod   7043 non-null    object 
 18  MonthlyCharges 7043 non-null    float64
 19  TotalCharges    7043 non-null    object 
 20  Churn           7043 non-null    object 
 21 dtype: object
```

```
0    customerID      7043 non-null   object
1    gender          7043 non-null   object
2    SeniorCitizen   7043 non-null   int64
3    Partner          7043 non-null   object
4    Dependents       7043 non-null   object
5    tenure           7043 non-null   int64
6    PhoneService     7043 non-null   object
7    MultipleLines    7043 non-null   object
8    InternetService  7043 non-null   object
9    OnlineSecurity   7043 non-null   object
10   OnlineBackup      7043 non-null   object
11   DeviceProtection 7043 non-null   object
12   TechSupport       7043 non-null   object
13   StreamingTV       7043 non-null   object
14   StreamingMovies   7043 non-null   object
15   Contract          7043 non-null   object
16   PaperlessBilling  7043 non-null   object
17   PaymentMethod     7043 non-null   object
18   MonthlyCharges    7043 non-null   float64
19   TotalCharges       7043 non-null   object
20   Churn             7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[120]: df.unique()
```

```
[120]: customerID      7043
gender          2
SeniorCitizen   2
Partner          2
Dependents       2
tenure           73
PhoneService     2
MultipleLines    3
InternetService  3
OnlineSecurity   3
OnlineBackup      3
DeviceProtection 3
TechSupport       3
StreamingTV       3
StreamingMovies   3
Contract          3
PaperlessBilling  2
PaymentMethod     4
MonthlyCharges    1585
TotalCharges      6531
Churn             2
```

```
dtype: int64
```

```
[121]: df.duplicated().sum()
```

```
[121]: np.int64(0)
```

```
[122]: df.describe()
```

```
[122]:   SeniorCitizen      tenure  MonthlyCharges
count    7043.000000  7043.000000    7043.000000
mean     0.162147    32.371149    64.761692
std      0.368612    24.559481    30.090047
min      0.000000    0.000000    18.250000
25%     0.000000    9.000000    35.500000
50%     0.000000   29.000000   70.350000
75%     0.000000   55.000000   89.850000
max     1.000000   72.000000  118.750000
```

```
[123]: # Convert 'TotalCharges' to numeric, coerce errors to NaN
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")
```

```
[124]: # Check missing values across all columns
missing_values = df.isnull().sum()
print("Missing values per column:\n", missing_values)
```

Missing values per column:

```
customerID          0
gender              0
SeniorCitizen       0
Partner             0
Dependents          0
tenure              0
PhoneService        0
MultipleLines        0
InternetService     0
OnlineSecurity      0
OnlineBackup         0
DeviceProtection    0
TechSupport          0
StreamingTV         0
StreamingMovies     0
Contract            0
PaperlessBilling    0
PaymentMethod        0
MonthlyCharges      0
TotalCharges        11
Churn               0
dtype: int64
```

```
[125]: # Optional: If any other numeric columns had missing values, fill with median
numeric_cols = df.select_dtypes(include=["int64", "float64"]).columns

for col in numeric_cols:
    if df[col].isnull().sum() > 0:
        median_val = df[col].median()
        df[col].fillna(median_val, inplace=True)
        print(f"Imputed missing values in {col} with median: {median_val}")
```

Imputed missing values in TotalCharges with median: 1397.475

/tmp/ipykernel_4624/900270614.py:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna(median_val, inplace=True)
```

```
[126]: # Check missing values across all columns
missing_values = df.isnull().sum()
print("Missing values per column:\n", missing_values)
```

Missing values per column:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0

```
TotalCharges      0  
Churn            0  
dtype: int64
```

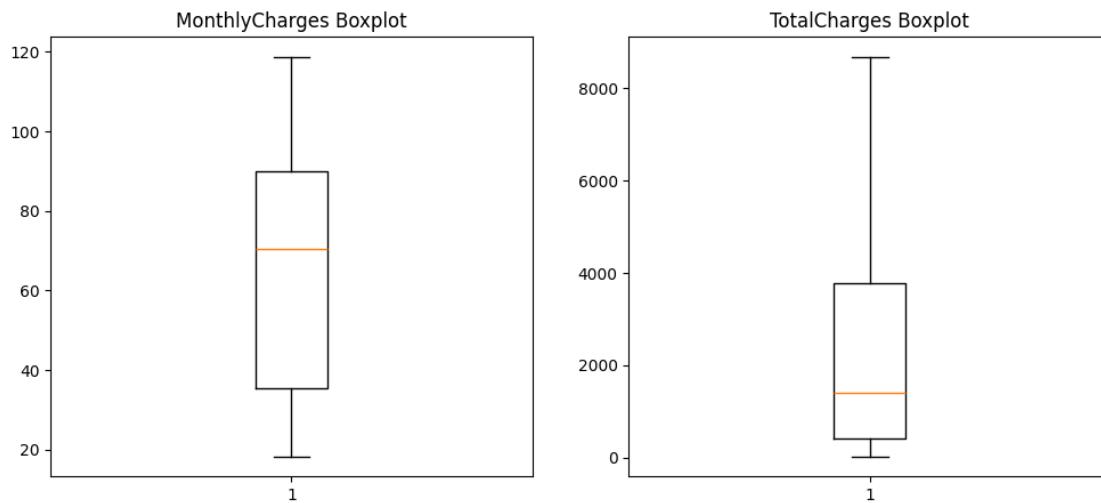
```
[127]: # step 2: outlier detection and treatment  
import pandas as pd  
import matplotlib.pyplot as plt  
  
# Step 1: Visualize outliers  
plt.figure(figsize=(12,5))  
  
plt.subplot(1, 2, 1)  
plt.boxplot(df['MonthlyCharges'])  
plt.title('MonthlyCharges Boxplot')  
  
plt.subplot(1, 2, 2)  
plt.boxplot(df['TotalCharges'])  
plt.title('TotalCharges Boxplot')  
  
plt.show()  
  
# Step 2: Detect outliers using IQR  
def detect_outliers_iqr(data, column):  
    Q1 = data[column].quantile(0.25)  
    Q3 = data[column].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]  
    return outliers  
  
monthly_outliers = detect_outliers_iqr(df, 'MonthlyCharges')  
total_outliers = detect_outliers_iqr(df, 'TotalCharges')  
  
print("MonthlyCharges outliers count:", monthly_outliers.shape[0])  
print("TotalCharges outliers count:", total_outliers.shape[0])  
  
# Step 3: Cap outliers (Winsorization)  
def cap_outliers_iqr(data, column):  
    Q1 = data[column].quantile(0.25)  
    Q3 = data[column].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
    data[column] = data[column].clip(lower=lower_bound, upper=upper_bound)  
  
cap_outliers_iqr(df, 'MonthlyCharges')
```

```

cap_outliers_iqr(df, 'TotalCharges')

# Step 4: Verify
print("After capping:")
print(df[['MonthlyCharges', 'TotalCharges']].describe())

```



```

MonthlyCharges outliers count: 0
TotalCharges outliers count: 0
After capping:
      MonthlyCharges  TotalCharges
count      7043.000000  7043.000000
mean       64.761692   2281.916928
std        30.090047   2265.270398
min        18.250000   18.800000
25%       35.500000   402.225000
50%       70.350000  1397.475000
75%       89.850000  3786.600000
max       118.750000  8684.800000

```

```

[142]: # -----
# Data Transformation: Categorical Encoding + Tenure Grouping
# ----- 

# 1 Transform 'tenure' into interpretable bins
df["TenureGroup"] = pd.cut(
    df["tenure"],
    bins=[0, 12, 24, 48, 72],
    labels=["0-1yr", "1-2yr", "2-4yr", "4-6yr"]
)

```

```

# 2 Identify categorical and numerical features
categorical_features = [
    "gender", "Partner", "Dependents", "PhoneService", "MultipleLines",
    "InternetService", "OnlineSecurity", "OnlineBackup", "DeviceProtection",
    "TechSupport", "StreamingTV", "StreamingMovies", "Contract",
    "PaperlessBilling", "PaymentMethod", "TenureGroup"
]

numerical_features = ["SeniorCitizen", "tenure", "MonthlyCharges"]

# 3 Build preprocessing pipelines
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median"))
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(drop="first", handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(transformers=[
    ("num", numeric_transformer, numerical_features),
    ("cat", categorical_transformer, categorical_features)
])

# 4 Apply transformations
X = df[numerical_features + categorical_features]
y = df["Churn"].map({"No": 0, "Yes": 1})

X_transformed = preprocessor.fit_transform(X)

print("Data Transformation completed!")
print("Transformed feature shape:", X_transformed.shape)

```

Data Transformation completed!
 Transformed feature shape: (7043, 32)

[145]: #Drop irrelevant columns
 df.drop(['customerID'], axis=1, inplace=True, errors='ignore')
 df.head()

```
[145]:    gender SeniorCitizen Partner Dependents tenure PhoneService \
0 Female 0 Yes No 1 No
1 Male 0 No No 34 Yes
2 Male 0 No No 2 Yes
3 Male 0 No No 45 No
4 Female 0 No No 2 Yes

    MultipleLines InternetService OnlineSecurity OnlineBackup ... \
0 No phone service DSL No Yes ...
1 No DSL Yes No ...
2 No DSL Yes Yes ...
3 No phone service DSL Yes No ...
4 No Fiber optic No No ...

    MonthlyCharges TotalCharges Churn TenureGroup Speed_x DataAllowance_x \
0 29.85 29.85 No 0-1yr 25.0 500.0
1 56.95 1889.50 No 2-4yr 25.0 500.0
2 53.85 108.15 Yes 0-1yr 25.0 500.0
3 42.30 1840.75 No 2-4yr 25.0 500.0
4 70.70 151.65 Yes 0-1yr 100.0 2000.0

    Speed_y DataAllowance_y Speed DataAllowance
0 25.0 500.0 25.0 500.0
1 25.0 500.0 25.0 500.0
2 25.0 500.0 25.0 500.0
3 25.0 500.0 25.0 500.0
4 100.0 2000.0 100.0 2000.0

[5 rows x 27 columns]
```

```
[144]: # -----
# Correct numeric and categorical features
# -----
numeric_features = ["SeniorCitizen", "tenure", "MonthlyCharges", "Speed", "DataAllowance"]
categorical_features = [
    "gender", "Partner", "Dependents", "PhoneService", "MultipleLines",
    "InternetService", "OnlineSecurity", "OnlineBackup", "DeviceProtection",
    "TechSupport", "StreamingTV", "StreamingMovies",
    "Contract", "PaperlessBilling", "PaymentMethod", "TenureGroup"
]
# -----
# ColumnTransformer
# -----
numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median"))
```

```

])
categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(drop="first", handle_unknown="ignore"))
])
preprocessor = ColumnTransformer(transformers=[
    ("num", numeric_transformer, numeric_features),
    ("cat", categorical_transformer, categorical_features)
])
# -----
# Full Pipeline with SMOTE + Logistic Regression
# -----
model = ImbPipeline(steps=[
    ("preprocessor", preprocessor),
    ("smote", SMOTE(random_state=42)),
    ("classifier", LogisticRegression(max_iter=2000, solver="lbfgs"))
])
# -----
# Feature matrix and target
# -----
X = df[numeric_features + categorical_features]
y = df["Churn"].map({"No": 0, "Yes": 1})
# -----
# Train/test split
# -----
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
# -----
# Train the model
# -----
model.fit(X_train, y_train)
print("Pipeline with integrated internet plans trained successfully!")

# -----
# Check transformed feature shape
# -----
X_transformed = preprocessor.fit_transform(X)
print("Transformed feature shape:", X_transformed.shape)

```

Pipeline with integrated internet plans trained successfully!

Transformed feature shape: (7043, 34)

```
[21]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import VarianceThreshold
from sklearn.linear_model import LogisticRegression
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
from sklearn.model_selection import train_test_split

# -----
# Define raw features and target
# -----
raw_features = [
    "gender", "SeniorCitizen", "Partner", "Dependents", "tenure",
    "PhoneService", "MultipleLines", "InternetService", "OnlineSecurity",
    "OnlineBackup", "DeviceProtection", "TechSupport", "StreamingTV",
    "StreamingMovies", "Contract", "PaperlessBilling", "PaymentMethod",
    "MonthlyCharges", "TenureGroup"
]

X = df[raw_features].copy()
y = df["Churn"].map({"No": 0, "Yes": 1})

# -----
# Identify numeric and categorical features
# -----
numeric_features = ["SeniorCitizen", "tenure", "MonthlyCharges"]
categorical_features = [
    "gender", "Partner", "Dependents", "PhoneService", "MultipleLines",
    "InternetService", "OnlineSecurity", "OnlineBackup", "DeviceProtection",
    "TechSupport", "StreamingTV", "StreamingMovies", "Contract",
    "PaperlessBilling", "PaymentMethod", "TenureGroup"
]

# -----
# Build preprocessing pipelines
# -----
numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median"))
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(drop="first", handle_unknown="ignore"))
```

```

])
```

```

preprocessor = ColumnTransformer(transformers=[  

    ("num", numeric_transformer, numeric_features),  

    ("cat", categorical_transformer, categorical_features)
])
```

```

# -----  

# Build full pipeline with VarianceThreshold (Data Reduction) + SMOTE +  

# Logistic Regression  

# -----
```

```

model = ImbPipeline(steps=[  

    ("preprocessor", preprocessor),  

    ("variance_threshold", VarianceThreshold(threshold=0.01)), # Data Reduction  

    ("smote", SMOTE(random_state=42)),  

    ("classifier", LogisticRegression(max_iter=2000, solver="lbfgs"))
])
```

```

# -----  

# Train/test split  

# -----
```

```

X_train, X_test, y_train, y_test = train_test_split(  

    X, y, test_size=0.2, stratify=y, random_state=42
)
```

```

# -----  

# Train the pipeline  

# -----
```

```

model.fit(X_train, y_train)  

print("Pipeline with Data Reduction trained successfully!")
```

```

# -----  

# Check which features were kept after reduction  

# -----
```

```

selected_features_mask = model.named_steps['variance_threshold'].get_support()  

preprocessed_features = model.named_steps['preprocessor'].  

    get_feature_names_out()  

selected_features = preprocessed_features[selected_features_mask]
```

```

print("Number of features before reduction:", len(preprocessed_features))  

print("Number of features after reduction:", len(selected_features))  

print("Selected Features:", selected_features)
```

Pipeline with Data Reduction trained successfully!
Number of features before reduction: 32
Number of features after reduction: 32
Selected Features: ['num_SeniorCitizen' 'num_tenure' 'num_MonthlyCharges'

```
'cat__gender_Male' 'cat__Partner_Yes' 'cat__Dependents_Yes'
'cat__PhoneService_Yes' 'cat__MultipleLines_No phone service'
'cat__MultipleLines_Yes' 'cat__InternetService_Fiber optic'
'cat__InternetService_No' 'cat__OnlineSecurity_No internet service'
'cat__OnlineSecurity_Yes' 'cat__OnlineBackup_No internet service'
'cat__OnlineBackup_Yes' 'cat__DeviceProtection_No internet service'
'cat__DeviceProtection_Yes' 'cat__TechSupport_No internet service'
'cat__TechSupport_Yes' 'cat__StreamingTV_No internet service'
'cat__StreamingTV_Yes' 'cat__StreamingMovies_No internet service'
'cat__StreamingMovies_Yes' 'cat__Contract_One year'
'cat__Contract_Two year' 'cat__PaperlessBilling_Yes'
'cat__PaymentMethod_Credit card (automatic)'
'cat__PaymentMethod_Electronic check' 'cat__PaymentMethod_Mailed check'
'cat__TenureGroup_1-2yr' 'cat__TenureGroup_2-4yr'
'cat__TenureGroup_4-6yr']
```

[150]: # Count number of churned and non-churned customers

```
churn_counts = df['Churn'].value_counts()
print("Churn class distribution:")
print(churn_counts)

# Calculate percentage of each class
churn_percent = df['Churn'].value_counts(normalize=True) * 100
print("\nChurn class distribution (%):")
print(churn_percent)
```

Churn class distribution:

```
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

Churn class distribution (%):

```
Churn
No      73.463013
Yes     26.536987
Name: proportion, dtype: float64
```

[148]: from sklearn.preprocessing import OneHotEncoder

```
from imblearn.over_sampling import SMOTE
import pandas as pd

# One-hot encode categorical columns
X_encoded = pd.get_dummies(X_train, drop_first=True)

# Apply SMOTE
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_encoded, y_train)
```

```
# Check new distribution
from collections import Counter
print("Training set class distribution after SMOTE:")
print(Counter(y_train_resampled))
```

Training set class distribution after SMOTE:
 Counter({0: 4139, 1: 4139})

[]:

[]: #ANSWERS

1. Handling Missing Values:

For handling missing values **in** the Telco dataset, especially **in** the columns **like "Monthly Charges" and "Total Charges"**, imputation techniques such **as mean, median, or mode** are commonly used. For **"Monthly Charges" and "Total Charges"**, imputation **with** the median **or** mean can work well. The median **is particularly useful if the distribution is skewed, while the mean might be more appropriate if the data is normally distributed.**

Effectively handling missing values **is** essential because features like **"Monthly Charges" and "Total Charges"** are key predictors of customer churn. Proper **imputation ensures valuable data **is** retained, allowing the model to leverage all available information. Failing to address missing values can lead to biased predictions and decreased model performance.**

2. Handling Outliers:

Outliers **in "Monthly Charges" or "Total Charges"** might be genuine (such **as high-paying customers**) **or** due to errors. Instead of removing them without investigation, it's best to first analyze whether the outliers are valid.

If they represent valid customer behavior, they should be kept **as** they could provide valuable insights **for** predicting churn. However, **if** the outliers are **errors or do not contribute meaningfully to the prediction, they should be removed or corrected**. This approach ensures that we maintain valuable data **while improving the model's accuracy**.

3. Handling Categorical Values:

For features like **"Contract", "Internet Service", and "Phone Service"**, we need **to convert these categorical variables into numerical values that the machine learning model can interpret**.

Label Encoding:

For ordinal categorical variables (where the values have a meaningful order), such as "Contract" with values like "Month-to-Month", "One Year", and "Two Year", label encoding can be used. Each category would be assigned an integer (e.g., "Month-to-Month" = 0, "One Year" = 1, "Two Year" = 2).

One-Hot Encoding:

For nominal categorical variables (where there is no meaningful order), such as "Internet Service" or "Phone Service", one-hot encoding should be used. This creates a binary column for each category (e.g., for "Internet Service", there would be columns for "Fiber Optic" and "DSL", and the values would be 0 or 1 depending on the customer's service type).

Proper encoding ensures the machine learning model understands the relationship between categorical variables and the target variable (churn). This transformation allows the model to work effectively with non-numerical data, making it possible to make accurate predictions.

4. Handling the Tenure Column:

Binning and Mapping:

We transformed the "Tenure" feature into categories using binning. We divided the tenure into four categories:

"Short" (0-12 months)

"Medium" (13-24 months)

"Long" (25-36 months)

"Very Long" (36+ months)

Each category was then mapped to a numerical value:

"Short" = 0

"Medium" = 1

"Long" = 2

"Very Long" = 3

Finally, we dropped the original "Tenure" and "TenureCategory" columns, leaving only the transformed numeric version for use in the model.

Binning and mapping "Tenure" into categories allows the model to capture patterns related to customer loyalty more effectively. It simplifies the data, making it easier for the model to identify trends, and prevents the model from being impacted by extreme tenure values. This transformation also improves interpretability by converting the feature into meaningful categories.

5. Data Integration:

We would have combined the service information (like "Speed" and "Data Allowance") with the Telco dataset by matching them using a shared column called "CustomerID".

After combining the data, we would have fixed the missing details for some customers. For "Speed", we would have filled the gaps with the most common value (mode). For "Data Allowance", we would have used the average (mean) to fill the gaps.

We would have also added a new column called "ServiceInfoMissing" to mark customers with missing details. If a customer was missing data, this column would have shown a value of 1; otherwise, it would have shown 0.

Adding the service information would have improved the dataset for predicting which customers might leave (churn). Fixing and marking missing details would have kept the data accurate and reliable for the prediction model.

Code Example for the Above

```
# Merging datasets on CustomerID
combined_data = pd.merge(data, service_data, on='CustomerID', how='left')

# 'service_data' is the data we originally have
# service_data is the additional data with service information

# Impute missing values
combined_data['Speed'].fillna(combined_data['Speed'].mode()[0], inplace=True)
combined_data['Data Allowance'].fillna(combined_data['Data Allowance'].mean(), inplace=True)

# Flag missing data
combined_data['ServiceInfoMissing'] = combined_data[['Speed', 'Data Allowance']].isnull().any(axis=1).astype(int)
```

Error Message:

NameError

Traceback (most recent call last)

```
Cell In[70], line 3
1 # Merging datasets on CustomerID
2 combined_data = pd.merge(data, service_data, on='CustomerID', how='left')
3
NameError: name 'service_data' is not defined
```

6. Feature Selection:

For feature selection, we used Recursive Feature Elimination (RFE) to identify and remove irrelevant features from the Telco dataset. RFE helps select the most important features by recursively removing the least important ones based on model performance. By applying RFE, we eliminated features like "CustomerID" that are not useful for churn prediction, improving model efficiency and focusing on the most relevant data.

7. Data Imbalance:

As stated, customer churn is a relatively rare event compared to non-churn, which can lead to class imbalance in the dataset. This imbalance may cause the model to be biased toward predicting the majority class (non-churn), potentially reducing the model's ability to identify churned customers accurately.

To address this, we applied SMOTEENN (Synthetic Minority Over-sampling Technique + Edited Nearest Neighbors). SMOTEENN is an effective technique for handling class imbalance because it works by generating synthetic samples for the minority class (churn) while removing noise from the dataset using Edited Nearest Neighbors.

This approach helps balance the dataset by generating samples for churned customers, which improves the model's ability to predict both churn and non-churn instances accurately.